

```
apiVersion: v1
kind: Pod
metadata:
  name: example-pod
spec:
  containers:
  - name: container
    image: nginx
    ports:
    - containerPort: 80
    livenessProbe:
      httpGet:
        path: /index.html
        port: 80
      initialDelaySeconds: 15
      periodSeconds: 15
    readinessProbe:
      httpGet:
        path: /index.html
        port: 80
      initialDelaySeconds: 5
      periodSeconds: 10
```

```
~
~
~
~
~
~
~
```

```

controlplane:~$ ls
filesystem liveness_readiness.yaml
controlplane:~$ vi liveness_readiness.yaml
controlplane:~$ kubectl apply -f liveness_readiness.yaml ✓
pod/example-pod created
controlplane:~$ kubectl get pods ✓
NAME          READY   STATUS             RESTARTS   AGE
example-pod   0/1     ContainerCreating   0           6s
controlplane:~$ kubectl get pods
NAME          READY   STATUS             RESTARTS   AGE
example-pod   0/1     Running            0           12s
controlplane:~$ kubectl get pods
NAME          READY   STATUS             RESTARTS   AGE
example-pod   0/1     Running            0           14s
controlplane:~$ kubectl get pods
NAME          READY   STATUS             RESTARTS   AGE
example-pod   0/1     Running            0           17s
controlplane:~$ kubectl get pods ✓
NAME          READY   STATUS             RESTARTS   AGE
example-pod   1/1     Running            0           25s
controlplane:~$ kubectl exec -it example-pod -- bin/bash ✓
root@example-pod:/# cls
bash: cls: command not found
root@example-pod:/# ls
bin  dev                docker-entrypoint.sh  home  lib64  mnt  proc  run  srv  tmp  var
boot  docker-entrypoint.d  etc                  lib   media  opt  root  sbin  sys  usr
root@example-pod:/# cd usr
root@example-pod:/usr# ls
bin  games  include  lib  lib64  libexec  local  sbin  share  src
root@example-pod:/usr# cd share
root@example-pod:/usr/share# ls
X11          bug          debianutils  dpkg        gdb          libc-bin    man          nginx        pixmaps     util-linux
base-files   ca-certificates  dict         fontconfig  info         libgcrypt20  maven-repo  pam          polkit-1    xml
base-passwd  common-licenses  doc          fonts       java         lintian     menu         pam-configs  tabset      zoneinfo
bash-completion  debconf        doc-base    gcc          keyrings     locale      misc         perl5        terminfo    zsh
root@example-pod:/usr/share# cd nginx
root@example-pod:/usr/share/nginx# ls
tml
root@example-pod:/usr/share/nginx# cd html
root@example-pod:/usr/share/nginx/html# ls
0x.html  index.html
root@example-pod:/usr/share/nginx/html# rm index.html ✓
root@example-pod:/usr/share/nginx/html# ls
0x.html
root@example-pod:/usr/share/nginx/html# exit ✓
xit
controlplane:~$ kubectl get pods
NAME          READY   STATUS             RESTARTS   AGE
example-pod   1/1     Running            0           2m15s
controlplane:~$ kubectl get pods
NAME          READY   STATUS             RESTARTS   AGE
example-pod   1/1     Running            0           2m19s
controlplane:~$ kubectl get pods
NAME          READY   STATUS             RESTARTS   AGE
example-pod   1/1     Running            0           2m26s
controlplane:~$ kubectl get pods ✓
NAME          READY   STATUS             RESTARTS   AGE
example-pod   0/1     Running            0           2m38s
controlplane:~$ kubectl get pods
NAME          READY   STATUS             RESTARTS   AGE
example-pod   0/1     Running            0           2m43s
controlplane:~$ kubectl get pods
NAME          READY   STATUS             RESTARTS   AGE
example-pod   0/1     Running            1 (9s ago)  2m54s
controlplane:~$ kubectl get pods
NAME          READY   STATUS             RESTARTS   AGE ✓
example-pod   1/1     Running            1 (15s ago)  3m
controlplane:~$ kubectl exec -it example-pod --bin/bash ✓
error: unknown flag: --bin/bash
see 'kubectl exec --help' for usage.

```

```
example-pod 1/1 Running 1 (15s ago) 3m
controlplane:~$ kubectl exec -it example-pod --bin/bash
error: unknown flag: --bin/bash
See 'kubectl exec --help' for usage.
controlplane:~$ kubectl exec -it example-pod -- bin/bash
root@example-pod:/# cd usr/share/nginx/html✓
root@example-pod:/usr/share/nginx/html# ls
50x.html index.html✓
root@example-pod:/usr/share/nginx/html# █
```

Horizontal Pod Autoscaling

```
Example Desktop Editor 1/1 Running 1 (15s ago) 3m
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: php-apache
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: nginx-deployment
  minReplicas: 1
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 50
~
~
```

```

apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io created
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl top pods
error: Metrics API not available
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ vi metrics-server.yaml
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl top pods
NAME          CPU(cores)   MEMORY(bytes)
example-pod    1m           2Mi
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl get nodes
NAME          STATUS    ROLES    AGE   VERSION
controlplane  Ready     control-plane  22d   v1.32.1
node01        Ready     <none>      22d   v1.32.1
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl top nodes
NAME          CPU(cores)   CPU(%)   MEMORY(bytes)   MEMORY(%)
controlplane  101m         10%      1134Mi          60%
node01        120m         12%      907Mi           48%
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ vi deploy.yaml
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl apply -f deploy.yaml
deployment.apps/nginx-deployment created
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
example-pod   1/1     Running   1 (27m ago) 30m
nginx-deployment-59cf44d54c-c4c4s  0/1     ContainerCreating  0           5s
nginx-deployment-59cf44d54c-n6cbt  0/1     ContainerCreating  0           5s
nginx-deployment-59cf44d54c-vnr8x  0/1     ContainerCreating  0           5s
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
example-pod   1/1     Running   1 (27m ago) 30m
nginx-deployment-59cf44d54c-c4c4s  1/1     Running   0           9s
nginx-deployment-59cf44d54c-n6cbt  1/1     Running   0           9s
nginx-deployment-59cf44d54c-vnr8x  1/1     Running   0           9s
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ vi hps.yaml
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ vi hps.yaml
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ kubcetl apply -f hps.yaml
kubcetl: command not found
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl apply -f hps.yaml
kubcetl: command not found
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl apply -f hps.yaml
horizontalpodautoscaler.autoscaling/php-apache created
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl get hpa
NAME          REFERENCE          TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
php-apache    Deployment/nginx-deployment  cpu: 0%/50%  1         10        3          20s
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
example-pod   1/1     Running   1 (38m ago) 40m
nginx-deployment-59cf44d54c-c4c4s  1/1     Running   0           10m
nginx-deployment-59cf44d54c-n6cbt  1/1     Running   0           10m
nginx-deployment-59cf44d54c-vnr8x  1/1     Running   0           10m
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl describe pod php-apache
Error from server (NotFound): pods "php-apache" not found
controlplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl describe hpa php-apache
Name:         php-apache
Namespace:    default
Labels:       <none>
Annotations:  <none>
CreationTimestamp: Mon, 14 Apr 2025 14:35:23 +0000
Reference:    Deployment/nginx-deployment
Metrics:      ( current / target )
  resource cpu on pods (as a percentage of request): 0% (0) / 50%
Min replicas: 1
Max replicas: 10
Deployment pods: 3 current / 3 desired
Conditions:
  Type           Status  Reason
  ----           -
  AbleToScale    True    ScaleDownStabilized recent recommendations were higher than current one, applying the highest recent recommendation
  ScalingActive  True    ValidMetricFound    the HPA was able to successfully calculate a replica count from cpu resource utilization (percent
  ge of request)
  ScalingLimited False   DesiredWithinRange  the desired count is within the acceptable range
Events:         <none>
controlplane:~/kubernetes-training/k8-manifests/metrics_server$

```

```

memory: "64Mi"
cpu: "250m"
limits:
  memory: "128Mi"
  cpu: "500m"

ontrolplane:~/kubernetes-training/k8-manifests/metrics_server$ ^C
ontrolplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
example-pod                         1/1     Running   1 (42m ago)  45m
nginx-deployment-59cf44d54c-vnr8x  1/1     Running   0           14m
ontrolplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl describe hpa php-apache
Name:                               php-apache
Namespace:                          default
Labels:                              <none>
Annotations:                         <none>
CreationTimestamp:                  Mon, 14 Apr 2025 14:35:23 +0000
Reference:                          Deployment/nginx-deployment
Metrics:                             ( current / target )
  resource cpu on pods  (as a percentage of request):  0% (0) / 50%
Min replicas:                               1
Max replicas:                               10
Deployment pods:                             1 current / 1 desired
Conditions:
  Type            Status  Reason                        Message
  ----            -
  AbleToScale     True    ReadyForNewScale             recommended size matches current size
  ScalingActive   True    ValidMetricFound             the HPA was able to successfully calculate a replica count from cpu resource utilization (percentage of request)
  ScalingLimited  True    TooFewReplicas               the desired replica count is less than the minimum replica count
Events:
  Type            Reason                  Age    From                      Message
  ----            -
  Normal          SuccessfulRescale       40s    horizontal-pod-autoscaler  New size: 1; reason: All metrics below target
ontrolplane:~/kubernetes-training/k8-manifests/metrics_server$

memory: "64Mi"
cpu: "250m"
limits:
  memory: "128Mi"
  cpu: "500m"

ontrolplane:~/kubernetes-training/k8-manifests/metrics_server$ ^C
ontrolplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
example-pod                         1/1     Running   1 (42m ago)  45m
nginx-deployment-59cf44d54c-vnr8x  1/1     Running   0           14m
ontrolplane:~/kubernetes-training/k8-manifests/metrics_server$ kubectl describe hpa php-apache
Name:                               php-apache
Namespace:                          default
Labels:                              <none>
Annotations:                         <none>
CreationTimestamp:                  Mon, 14 Apr 2025 14:35:23 +0000
Reference:                          Deployment/nginx-deployment
Metrics:                             ( current / target )
  resource cpu on pods  (as a percentage of request):  0% (0) / 50%
Min replicas:                               1
Max replicas:                               10
Deployment pods:                             1 current / 1 desired
Conditions:
  Type            Status  Reason                        Message
  ----            -
  AbleToScale     True    ReadyForNewScale             recommended size matches current size
  ScalingActive   True    ValidMetricFound             the HPA was able to successfully calculate a replica count from cpu resource utilization (percentage of request)
  ScalingLimited  True    TooFewReplicas               the desired replica count is less than the minimum replica count
Events:
  Type            Reason                  Age    From                      Message
  ----            -
  Normal          SuccessfulRescale       40s    horizontal-pod-autoscaler  New size: 1; reason: All metrics below target
ontrolplane:~/kubernetes-training/k8-manifests/metrics_server$

```