

Node Selector

```
ubuntu@k8-master:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
k8-master            Ready     control-plane  24h   v1.31.7
k8-worker1           Ready     <none>      24h   v1.31.7
k8-worker2           Ready     <none>      24h   v1.31.7
ubuntu@k8-master:~$ kubectl get nodes --show-labels
NAME                STATUS    ROLES    AGE   VERSION   LABELS
k8-master            Ready     control-plane  24h   v1.31.7   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=k8-master,kubernetes.io/
linux,node-role.kubernetes.io/control-plane=,node.kubernetes.io/exclude-from-external-load-balancers=
k8-worker1           Ready     <none>      24h   v1.31.7   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=k8-worker1,kubernetes.io/
linux
k8-worker2           Ready     <none>      24h   v1.31.7   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=k8-worker2,kubernetes.io/
linux

ubuntu@k8-master:~$ kubectl get nodes --show-labels
NAME                STATUS    ROLES    AGE   VERSION   LABELS
k8-master            Ready     control-plane  24h   v1.31.7   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=k8-master,kubernetes.io/
linux,node-role.kubernetes.io/control-plane=,node.kubernetes.io/exclude-from-external-load-balancers=
k8-worker1           Ready     <none>      24h   v1.31.7   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=k8-worker1,kubernetes.io/
linux
k8-worker2           Ready     <none>      24h   v1.31.7   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=k8-worker2,kubernetes.io/
linux

ubuntu@k8-master:~$ kubectl label nodes k8-worker1 cup=two
node/k8-worker1 labeled
ubuntu@k8-master:~$ kubectl get nodes --show-labels
NAME                STATUS    ROLES    AGE   VERSION   LABELS
k8-master            Ready     control-plane  24h   v1.31.7   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=k8-master,kubernetes.io/
linux,node-role.kubernetes.io/control-plane=,node.kubernetes.io/exclude-from-external-load-balancers=
k8-worker1           Ready     <none>      24h   v1.31.7   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,cup=two,kubernetes.io/arch=amd64,kubernetes.io/hostname=k8-worker1,kubernetes
.io/os=linux
k8-worker2           Ready     <none>      24h   v1.31.7   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=k8-worker2,kubernetes.io/
linux

ubuntu@k8-master:~$ kubectl label nodes k8-worker1 cup=two-
error: invalid label value: "cup=two-": a valid label must be an empty string or consist of alphanumeric characters, '-', '_' or '.', and must start and end with an alphanumeric character (
e.g. 'MyValue', or 'my_value', or '12345', regex used for validation is '(([A-Za-z0-9]([-A-Za-z0-9_\.])*)?([A-Za-z0-9])?)')
ubuntu@k8-master:~$ kubectl label nodes k8-worker1 cup-
node/k8-worker1 unlabeled
ubuntu@k8-master:~$ kubectl label nodes k8-worker1 cup=two
node/k8-worker1 labeled
ubuntu@k8-master:~$ kubectl get nodes --show-labels
NAME                STATUS    ROLES    AGE   VERSION   LABELS
k8-master            Ready     control-plane  24h   v1.31.7   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=k8-master,kubernetes.io/
linux,node-role.kubernetes.io/control-plane=,node.kubernetes.io/exclude-from-external-load-balancers=
k8-worker1           Ready     <none>      24h   v1.31.7   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,cup=two,kubernetes.io/arch=amd64,kubernetes.io/hostname=k8-worker1,kubernetes
.io/os=linux
k8-worker2           Ready     <none>      24h   v1.31.7   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=k8-worker2,kubernetes.io/
linux
```

I attached a label to the k8-worker1 node using the kubectl label command and verified whether the label was attached using kubectl get nodes --show-labels.

```
ubuntu@k8-master:~$ cat nodeselector.yaml
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    env: test
spec:
  containers:
  - name: nginx
    image: nginx
    imagePullPolicy: IfNotPresent
  nodeSelector:
    cpu: two
```

Created a nodeselector.yaml file using pod and give nodeSelector to it.

```

buntu@k8-master:~$ vi nodeselector.yaml
buntu@k8-master:~$ kubectl apply -f nodeselector.yaml
pod/nginx created
buntu@k8-master:~$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
nginx                               1/1     Running   0          6s
nginx-7fcb78bd77-7fkgb             1/1     Running   1 (4m56s ago)    21h
nginx-7fcb78bd77-c9wg8             1/1     Running   1 (4m56s ago)    21h
nginx-7fcb78bd77-m8rcg             1/1     Running   1 (4m53s ago)    21h
nginx1-965647b5c-6sg5b             1/1     Running   1 (4m56s ago)    20h
nginx1-965647b5c-99x2m             1/1     Running   1 (4m53s ago)    20h
nginx1-965647b5c-q6c48             1/1     Running   1 (4m53s ago)    20h
buntu@k8-master:~$ kubectl get deploy
NAME    READY   UP-TO-DATE   AVAILABLE   AGE
nginx   3/3     3            3           21h
nginx1  3/3     3            3           20h
buntu@k8-master:~$ kubectl delete deploy nginx nginx1
deployment.apps "nginx" deleted
deployment.apps "nginx1" deleted
buntu@k8-master:~$ kubectl get pods
NAME    READY   STATUS    RESTARTS   AGE
nginx   1/1     Running   0          50s
buntu@k8-master:~$ kubectl get pods -o wide
NAME    READY   STATUS    RESTARTS   AGE   IP            NODE           NOMINATED NODE   READINESS GATES
nginx   1/1     Running   0          59s   10.44.0.4     k8-worker1     <none>           <none>
buntu@k8-master:~$ cat nodeselector.yaml

```

I applied the nodeselector.yaml file using the kubectl apply command. This will schedule the Pod on the node with matching labels specified in the NodeSelector.

taint and tolerations

```

ubuntu@k8-master:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
k8-master           Ready     control-plane   24h   v1.31.7
k8-worker1          Ready     <none>         24h   v1.31.7
k8-worker2          Ready     <none>         24h   v1.31.7
ubuntu@k8-master:~$ kubectl taint nodes k8-worker1 key1=value1:NoSchedule
node/k8-worker1 tainted
ubuntu@k8-master:~$ kubectl describe nodes k8-worker1 | grep -i taint
Taints:          key1=value1:NoSchedule
ubuntu@k8-master:~$ vi taint_tola.yaml

```

I applied a taint on the node k8-worker1 using the command kubectl taint node k8-worker1 key1=value1:NoSchedule. To verify that the taint was applied, I used kubectl describe node k8-worker1

```

ubuntu@k8-master:~$ kubectl get pods -o wide
No resources found in default namespace.
ubuntu@k8-master:~$ kubectl get pods
No resources found in default namespace.
ubuntu@k8-master:~$ kubectl get pods
No resources found in default namespace.
ubuntu@k8-master:~$ vi taint_tola.yaml
ubuntu@k8-master:~$ kubectl apply -f taint_tola.yaml
pod/nginx4 created
ubuntu@k8-master:~$ kubectl get pods -o wide
NAME    READY   STATUS    RESTARTS   AGE   IP            NODE           NOMINATED NODE   READINESS GATES
nginx4   1/1     Running   0          5s    10.36.0.1     k8-worker2     <none>           <none>
ubuntu@k8-master:~$ vi taint_tola.yaml
ubuntu@k8-master:~$ kubectl apply -f taint_tola.yaml
pod/nginx3 created
ubuntu@k8-master:~$ kubectl get pods -o wide
NAME    READY   STATUS    RESTARTS   AGE   IP            NODE           NOMINATED NODE   READINESS GATES
nginx3   1/1     Running   0          5s    10.44.0.1     k8-worker1     <none>           <none>
nginx4   1/1     Running   0          46s   10.36.0.1     k8-worker2     <none>           <none>
ubuntu@k8-master:~$ cat taint_tola.yaml

```

Created a Pod definition YAML file (taint_tola.yaml) with a toleration matching the taint. Pods without this toleration won't be scheduled on the node with the taint. However, Pods

with the toleration can be scheduled on any node, including those without the taint

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-notoleration
  labels:
    env: test
spec:
  containers:
  - name: nginx
    image: nginx
    imagePullPolicy: IfNotPresent
```

```
ubuntu@k8s-master:~$ vi no_tola.yaml
ubuntu@k8s-master:~$ vi no_tola.yaml
ubuntu@k8s-master:~$ kubectl apply -f no_tola.yaml ✓
pod/nginx-notoleration created
ubuntu@k8s-master:~$ kubectl get pods -o wide ✓
NAME          READY   STATUS    RESTARTS   AGE   IP            NODE          NOMINATED NODE   READINESS GATES
nginx-notoleration 1/1     Running   0           10s   10.36.0.2     k8-worker2 ✓   <none>           <none>
nginx3         1/1     Running   0           3m53s 10.44.0.1     k8-worker1    <none>           <none>
nginx4         1/1     Running   0           4m34s 10.36.0.1     k8-worker2    <none>           <none>
ubuntu@k8s-master:~$ vi no_tola.yaml
ubuntu@k8s-master:~$ vi no_tola.yaml
ubuntu@k8s-master:~$ kubectl apply -f no_tola.yaml
pod/nginx-notoleration1 created
ubuntu@k8s-master:~$ kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP            NODE          NOMINATED NODE   READINESS GATES
nginx-notoleration 1/1     Running   0           78s   10.36.0.2     k8-worker2 ✓   <none>           <none>
nginx-notoleration1 1/1     Running   0           6s    10.36.0.3     k8-worker2 ✓   <none>           <none>
nginx3         1/1     Running   0           5m1s  10.44.0.1     k8-worker1    <none>           <none>
nginx4         1/1     Running   0           5m42s 10.36.0.1     k8-worker2    <none>           <none>
ubuntu@k8s-master:~$
```

Created a yaml file using pod definition without toleration. Now this pod cannot schedule on the node which have taint.

```

ubuntu@k8-master:~$ kubectl taint nodes k8-worker1 key1=value1:NoSchedule-
node/k8-worker1 untainted
ubuntu@k8-master:~$ kubectl taint nodes k8-worker1 key1=value1:PreferNoSchedule
node/k8-worker1 tainted
ubuntu@k8-master:~$ kubectl describe nodes k8-worker1 | grep -i taint
Taints:             key1=value1:PreferNoSchedule
ubuntu@k8-master:~$ vi no_tola.yaml
ubuntu@k8-master:~$ kubectl get pods -o wide
NAME                READY   STATUS    RESTARTS   AGE   IP            NODE       NOMINATED NODE   READINESS GATES
nginx-notoleration  1/1     Running   0           6m8s  10.36.0.2     k8-worker2 <none>          <none>
nginx-notoleration1 1/1     Running   0           4m56s 10.36.0.3     k8-worker2 <none>          <none>
nginx3               1/1     Running   0           9m51s 10.44.0.1     k8-worker1 <none>          <none>
nginx4               1/1     Running   0           10m    10.36.0.1     k8-worker2 <none>          <none>
ubuntu@k8-master:~$ kubectl apply -f no_tola.yaml
The Pod "nginx-PreferNoSchedule1" is invalid: metadata.name: Invalid value: "nginx-PreferNoSchedule1": a lowercase RFC 1123 subdomain must consist of lower case alphanumeric characters, '-' or '.', and must start and end with an alphanumeric character (e.g. 'example.com', regex used for validation is '[a-z0-9]([-a-z0-9.]?[a-z0-9])?')
ubuntu@k8-master:~$ vi no_tola.yaml
ubuntu@k8-master:~$ kubectl apply -f no_tola.yaml
pod/nginx-prefernoschedule1 created
ubuntu@k8-master:~$ kubectl get pods -o wide
NAME                READY   STATUS    RESTARTS   AGE   IP            NODE       NOMINATED NODE   READINESS GATES
nginx-notoleration  1/1     Running   0           8m47s 10.36.0.2     k8-worker2 <none>          <none>
nginx-notoleration1 1/1     Running   0           7m35s 10.36.0.3     k8-worker2 <none>          <none>
nginx-prefernoschedule1 1/1     Running   0           5s     10.36.0.4     k8-worker2 <none>          <none>
nginx3               1/1     Running   0           12m    10.44.0.1     k8-worker1 <none>          <none>
nginx4               1/1     Running   0           13m    10.36.0.1     k8-worker2 <none>          <none>
ubuntu@k8-master:~$ vi no_tola.yaml
ubuntu@k8-master:~$ kubectl apply -f no_tola.yaml
pod/nginx-prefernoschedule2 created
ubuntu@k8-master:~$ kubectl get pods -o wide
NAME                READY   STATUS    RESTARTS   AGE   IP            NODE       NOMINATED NODE   READINESS GATES
nginx-notoleration  1/1     Running   0           9m31s 10.36.0.2     k8-worker2 <none>          <none>
nginx-notoleration1 1/1     Running   0           8m19s 10.36.0.3     k8-worker2 <none>          <none>
nginx-prefernoschedule1 1/1     Running   0           49s    10.36.0.4     k8-worker2 <none>          <none>
nginx-prefernoschedule2 1/1     Running   0           10s    10.36.0.5     k8-worker2 <none>          <none>
nginx3               1/1     Running   0           13m    10.44.0.1     k8-worker1 <none>          <none>
nginx4               1/1     Running   0           13m    10.36.0.1     k8-worker2 <none>          <none>

```

Applied a taint with PreferNoSchedule to the node. This ensures that the scheduler will prefer not to schedule pods without a toleration, but if there are no other suitable nodes, the pod may still be scheduled on the tainted node

```

ubuntu@k8-master:~$ kubectl taint nodes k8-worker1 key1=value1:PreferNoSchedule-
node/k8-worker1 untainted
ubuntu@k8-master:~$ kubectl taint nodes k8-worker2 key1=value1:NoExecute
node/k8-worker2 tainted
ubuntu@k8-master:~$ kubectl get pods -o wide
NAME    READY   STATUS    RESTARTS   AGE   IP            NODE       NOMINATED NODE   READINESS GATES
nginx3  1/1     Running   0           15m    10.44.0.1     k8-worker1 <none>          <none>
ubuntu@k8-master:~$ kubectl get pods -o wide
NAME    READY   STATUS    RESTARTS   AGE   IP            NODE       NOMINATED NODE   READINESS GATES
nginx3  1/1     Running   0           15m    10.44.0.1     k8-worker1 <none>          <none>
ubuntu@k8-master:~$

```

Applied a taint with NoExecute that ensures no new pods will be scheduled on the node, and existing pods will be evicted (terminated) if they don't tolerate the taint